

Initial Experiments: Extending the MVP Simulator

William Dennis

June 2026

Extends the MVP simulator (`docs/mvp/mvp_specification.tex`) by enriching the state space, action space, and reward function, moving toward the full MDP defined in `docs/design/initial_design.tex`.

1 State Space Extension

References `docs/design/initial_design.tex` §3 and Issue #124.

1.1 Enriched StateView

1.2 Contextual Bandit Conditioning

1.3 Open Questions

2 Action Space Extension

References `docs/design/initial_design.tex` §3 and Issue #126.

2.1 ActionConfig

2.2 Transition Model Updates

2.3 Agent Compatibility

2.4 Open Questions

3 Reward Function Extension

References `docs/design/initial_design.tex` §3 and Issue #132.

3.1 Multi-Timescale Reward

3.2 Reward Composition

3.3 Synthetic Clinical Data

3.4 Open Questions

4 Combined Results